

Stochastic Anomaly Simulator

Project Documentation

Pole Anomaly Project

May 1, 2026

Contents

1	Introduction	2
2	Quick Setup Guide	3
2.1	Environment Setup	3
2.2	Dependencies	3
3	Synthetic Data Generation Pipeline (data-gen)	4
3.1	File Structure Review	4
3.2	Explanation of the Engine and Scripts	4
3.2.1	data-gen/src/anomaly_generator.py	4
3.2.2	data-gen/script/generate_custom_dataset.py	4
3.2.3	data-gen/script/test_generator_deformation.py	5
3.2.4	data-gen/script/visualize_anomalies_templates.py	5
3.3	Anomalies Templates Creation	5
3.3.1	Configuration Blueprint (CSV Logic)	5
4	Proposed Architectures for Anomaly Identification	8
4.1	General Data Preprocessing Pipeline	8
4.2	1D Convolutional Neural Networks (1D-CNN) / ResNet-1D	8
4.3	Convolutional Recurrent Neural Networks (CRNN)	9
4.4	Transformers with Multi-Head Self-Attention	9
4.5	1D U-Net (Segmentation Approach)	10
5	Reference PyTorch Implementation: 1D-CNN	11
5.1	Data Processing (dataset_loader.py & Generator)	11
5.2	Architecture Definition (model.py)	11
5.3	Training Loop (train.py)	11
5.4	Inference Visualization (visualize_inference.py)	12

1 Introduction

This document includes a general overview of the Pole Anomaly Project. The overall objective of this project is anomaly detection and classification using Neural Networks.

A critical component of this machine learning pipeline is the generation of training data. The data generation module is designed to inject, user-defined anomalies (such as Squares, Exponential Curves, and Bell Curves) into synthetic radiation noise based on real-world sensor baselines. The synthetic datasets generated from these anomaly templates will be directly used to train, validate, and test the neural networks for anomaly detection.

2 Quick Setup Guide

This project isolates dependencies using a Python virtual environment to ensure absolute reproducibility.

2.1 Environment Setup

To initialize the project on a new machine, execute the following commands in the root of the project:

```
# 1. Create a Python Virtual Environment
python -m venv .venv

# 2. Activate the Environment (Windows PowerShell)
.\.venv\Scripts\Activate.ps1

# 3. Install strictly required numerical libraries
pip install -r requirements.txt
```

2.2 Dependencies

The `requirements.txt` installs the following core standard math libraries:

- `numpy` & `pandas`: For fast mathematical array manipulations and template CSV handling.
- `plotly`: For high-fidelity, interactive visualizations saved in HTML logs.
- `scipy`, `statsmodels`, `tslearn`: Required specific dependencies for statistical interpolation and historical legacy dataset tests.

3 Synthetic Data Generation Pipeline (data-gen)

As the repository expands to include other projects (such as neural networks for anomaly detection), this section specifically details the anomaly generation pipeline.

3.1 File Structure Review

The repository enforces a strict separation of concerns, divided into modular directories:

- **data/**: Located at the root, contains raw background sensor readings (e.g., `2015_months_DebitDo`) used as a baseline.
- **data-gen/**: The main pipeline directory containing the synthetic dataset generator and its components:
 - **data-gen/anomalies/**: Contains the normalized `.csv` blueprints (templates) for different anomaly shapes (e.g., `exp_square.csv`, `bell.csv`, `m_sq.csv`).
 - **data-gen/data/**: The destination for the final generated output datasets (e.g., `synthetic_custom_dataset.csv`) utilized for neural network training.
 - **data-gen/logs/**: The destination for Plotly's interactable HTML visualizations and PNG exports. Provides a pure geometric proof of concept without requiring external dashboard readers.
 - **data-gen/src/**: Stores the pure mathematics and geometric generator libraries, decoupled from direct script logic.
 - **data-gen/script/**: Stores execution pipelines and visualization launchers. Scripts utilize the math from `data-gen/src/` to generate datasets or debug plots.

3.2 Explanation of the Engine and Scripts

3.2.1 data-gen/src/anomaly_generator.py

This is the core physics computational engine. It provides the `generate_anomaly()` function which transforms a normalized CSV blueprint into a physical discrete array anomaly string. The engine mathematically scales the template temporally (period span) and vertically (amplitude intensity) while embedding Gaussian "deformation" limits declared in the CSV file logic.

3.2.2 data-gen/script/generate_custom_dataset.py

The final execution pipeline for training. It parses the real noise characteristics (Mean, Standard Deviation) from the original physical gamma sensor. Using a baseline simulation of 10,000 to 50,000 steps, it spontaneously samples anomaly templates from the library, calculates randomized injection amplitudes (e.g., 5x to 25x STD multipliers) and duration bounds, and seamlessly injects them into the dataset, tagging absolute ground-truth labels alongside cleanly tracking the pure anomaly wave.




Figure 1: Output of `generate_custom_dataset.py`

3.2.3 `data-gen/script/test_generator_deformation.py`

A visual sandbox and boundary debugging tool. Iterating automatically over all discovered templates in the `data-gen/anomalies/` directory, it runs exactly 5 localized chaotic generation iterations of each template to visually verify and validate the mathematical variability limits enforced by the template's parameters. Exports directly to `anomaly_template_deformation_test.html` and `.png`.

3.2.4 `data-gen/script/visualize_anomalies_templates.py`

This script simply visualizes the pure original mathematical blueprints. It renders the pure lines and graphs the interpolation properties exactly as described in the blueprint CSV to a localized HTML and PNG, without injecting randomization or temporal noise.

3.3 Anomalies Templates Creation

Instead of relying on rigid pre-recorded dataset distributions like legacy UCR databases, targeted geometric structures are freely drawn and mathematically described in the `data-gen/anomalies/` folder.

3.3.1 Configuration Blueprint (CSV Logic)

A standard anomaly template requires 7 strict column definitions mapping logic coordinates and behavioral constraints:

`time, value, t_minus, t_plus, v_minus, v_plus, interp`

- **time & value:** Standardized physical shape coordinates (normally scaling from 0.0 to 1.0). They map the major mathematical inflection milestones of the curve.
- **Deformation Multipliers (`t_minus`, `t_plus`, `v_minus`, `v_plus`):** Defines spatial jitter limits per node. A value of 1.0 allows symmetric Gaussian distortion




Figure 2: Output of visualize_anomalies_templates.py




Figure 3: Output of test_generator_deformation.py

up to 100% of the global script variance parameter. A value of 0.0 creates an unyielding geometric wall perfectly locking the point in that specified space direction.

- **interp (Interpolation Mode):** Defines the localized mathematical connection trajectory bridging the gap toward the targeted point.
 - **linear:** Connects sequential milestones with a simple Euclidean linear straight interpolation fallback.
 - **exp:** Instigates a smooth sluggish temporal acceleration morphing into a harsh geometric peak, useful for raw exponential growth burst characterizations.
 - **bell:** Executes a fractional phase-shift cosine envelope ($\cos(\pi x)$), generating an intensely organic smooth "S-curve". Ensures a perfect non-rigid mathematical slope shift suitable for defining Gaussian-like topographic bell shapes.

4 Proposed Architectures for Anomaly Identification

Based on the characteristics of the synthetic dataset and the overall project objectives (classification performance, computational cost, and explainability), the following neural network architectures are proposed for anomaly detection and classification.

The simulated data presents specific challenges: extreme high-frequency gamma noise, distinct geometric signatures for anomalies (e.g., Bell, M, Square), and significant temporal and amplitude deformation. The models must be robust to background fluctuations and possess translation and scale invariance to generalize effectively.

4.1 General Data Preprocessing Pipeline

Regardless of the architecture chosen, the continuous sensor data cannot be processed infinitely. The shared preprocessing pipeline involves:

1. **Sliding Window Extraction:** Slicing the continuous gamma dose column into fixed-size overlapping windows (e.g., chunks of 256 or 512 time steps). This converts a single 1D time series of length N into a localized dataset of shape (**Samples**, **Window_Size**, **Channels**).
2. **Labeling Strategy:** For standard classification, assigning a categorical label to the whole window (e.g., 1 if an anomaly is present, 0 otherwise). For segmentation, assigning a binary label to every single point in the window to map the exact width of the anomaly.
3. **Normalization:** Applying Z-score normalization (subtracting the mean and dividing by the standard deviation) to each window independently. This ensures stability during gradient descent, as the extreme simulated radiation values can easily saturate activation functions like ReLU or Sigmoid.

4.2 1D Convolutional Neural Networks (1D-CNN) / ResNet-1D

1D convolutions are excellent for extracting local morphological features (like the sharp edges of a "Square" or the smooth curve of a "Bell") regardless of where they appear in the time window.

- **Performance & Cost:** Highly efficient and parallelizable, making them computationally cheap to train and deploy for real-time inference on continuous sensor data streams.
- **Explainability:** This architecture is directly compatible with Gradient-weighted Class Activation Mapping (**Grad-CAM**). Heatmaps can be generated over the input time series to visually prove that the model is activating on the structural milestones of the anomaly rather than simply overfitting to background noise.

Detailed Implementation Pipeline

1. **Input Layer:** Accepts a tensor of shape (**Batch_Size**, **Channels**, **Window_Size**) (in PyTorch) or (**Batch_Size**, **Window_Size**, **Channels**) (in Keras).

2. **Feature Extraction (Convolutional Blocks):** Stack 3 to 5 `Conv1D` layers. Each convolution slides a kernel over the time dimension to detect specific features. Follow each convolution with `BatchNormalization` to stabilize the learning process across batches, a `ReLU` activation to introduce necessary non-linearity, and `MaxPooling1D` to downsample the sequence length. This spatial reduction forces the network to learn higher-level features while increasing the number of filters (e.g., $32 \rightarrow 64 \rightarrow 128$).
3. **Global Pooling:** Use `GlobalAveragePooling1D` (or `AdaptiveAvgPool1d`) to collapse the time dimension entirely. This operation calculates the average output of each feature map across the time sequence, drastically reducing parameters, preventing overfitting, and acting as the structural prerequisite for Grad-CAM visualization.
4. **Output Layer:** A `Dense` (or `Linear`) layer with a `Softmax` activation for multi-class prediction, or a `Sigmoid` activation for binary classification.

4.3 Convolutional Recurrent Neural Networks (CRNN)

This hybrid approach utilizes a 1D-CNN base to extract local features, followed by a recurrent layer (such as an LSTM or GRU) to process the temporal sequence of those features. This is particularly useful for complex shapes like "M" where the model needs to track a structured sequence of events over time.

- **Performance & Cost:** Yields high accuracy for temporally stretched anomalies. The sequential processing of the RNN component adds a moderate computational overhead compared to a pure CNN.

Detailed Implementation Pipeline

1. **Spatial Feature Extractor:** Pass the input window through a shallow 1D-CNN (e.g., 2 `Conv1D` layers) *without* Global Pooling. The output remains a compressed sequence of dense feature vectors, retaining temporal order but with lower dimensionality.
2. **Temporal Sequence Modeler:** Pass the output sequence directly into a `Bidirectional` (LSTM) or GRU layer. The bidirectional nature allows the network to understand context from both the past and the future steps within the window, which is crucial for identifying shapes that have symmetrical properties (like Bells).
3. **Aggregation & Output Layer:** Extract the final hidden state of the RNN (or use an Attention layer over all hidden states) and pass it into a `Dense` layer for final classification.

4.4 Transformers with Multi-Head Self-Attention

Time-series Transformers excel at finding correlations across different parts of the signal. The multi-head self-attention mechanism can learn to suppress the high-frequency gamma noise and focus entirely on the broader structural patterns of the anomalies.

- **Explainability:** Attention mechanisms offer intrinsic explainability. The attention weights can be extracted and visualized to show exactly which time steps the model prioritized when making its classification decision, acting as a powerful and transparent complement to Grad-CAM.

Detailed Implementation Pipeline

1. **Patching/Embedding:** Unlike RNNs or CNNs, Transformers process all inputs simultaneously. Project the input sequence into a higher-dimensional embedding space (e.g., size 64) using a linear layer or a small `Conv1D` to create temporal "tokens".
2. **Positional Encoding:** Add explicit Sine/Cosine positional encodings to the embeddings so the model retains the chronological order of the time steps, since the attention mechanism itself is permutation-invariant.
3. **Transformer Blocks:** Pass the sequence through 2-4 `TransformerEncoder` blocks. Each block contains a Multi-Head Self-Attention mechanism (allowing tokens to interact and weigh each other's importance) and a Feed-Forward Network.
4. **Output Layer:** Aggregate the output via a dedicated CLS token (Class Token) or Global Average Pooling, and classify via a `Linear` layer.

4.5 1D U-Net (Segmentation Approach)

If the objective requires detecting the anomaly exactly at its onset by pinpointing the precise start and end timestamps, framing the problem as a time-series segmentation task is highly effective. A 1D U-Net or Fully Convolutional Network (FCN) outputs a probability score for every single time step, effectively generating a localized bounding mask over the anomaly in real-time.

Detailed Implementation Pipeline

1. **Encoder (Downsampling Pathway):** Use sequential `Conv1D` and `MaxPooling1D` layers to compress the time series. Crucially, save the output tensor at each resolution level before pooling. These are the "skip connections".
2. **Decoder (Upsampling Pathway):** Use `UpSampling1D` (or `Conv1DTranspose`) to gradually expand the sequence back to its original length. At each expansion step, concatenate the corresponding saved skip connection from the Encoder. This mechanism forces the network to combine deep semantic features (from the bottom of the U-Net) with high-resolution spatial context (from the skip connections).
3. **Point-wise Output Layer:** A `Conv1D` layer with a kernel size of 1 and exactly 1 filter per target class, followed by a Sigmoid/Softmax activation. The output shape perfectly matches the original input sequence length.
4. **Loss Function:** Utilize specialized loss functions like **Dice Loss** or **Focal Loss** to handle the severe class imbalance, as the vast majority of points in the dataset represent normal background noise.

5 Reference PyTorch Implementation: 1D-CNN

To validate the proposed architectures, a functional baseline implementation of the 1D-CNN was developed using PyTorch. The architecture has been specifically designed for **Multi-Class Classification**, capable of distinguishing between various anomaly templates (Bell, Square, M, etc.). The codebase is organized into modular scripts within the 1D-CNN directory:

5.1 Data Processing (dataset_loader.py & Generator)

The data generation pipeline was upgraded to assign specific integer IDs to each anomaly template, exporting a `class_mapping.txt` reference file. The `dataset_loader.py` module bridges this synthetic `.csv` generator and the neural network. It executes the sliding window extraction (default size 256, stride 32). It determines the window's label by extracting the maximum class ID present within the window, converting the sequence into a categorical multi-class target. It applies a `StandardScaler` and reshapes the tensors into the strict (`Batch_Size`, `Channels`, `Sequence_Length`) format required by PyTorch.

5.2 Architecture Definition (model.py)

The `Anomaly1DCNN` class inherits from `torch.nn.Module` and implements the exact 3-block convolutional architecture proposed:

- **Block 1:** `Conv1d(in=1, out=32, kernel=5) → BatchNorm1d → ReLU → MaxPool1d`.
- **Block 2:** `Conv1d(in=32, out=64, kernel=3) → BatchNorm1d → ReLU → MaxPool1d`.
- **Block 3:** `Conv1d(in=64, out=128, kernel=3) → BatchNorm1d → ReLU → MaxPool1d`.

The network utilizes an `AdaptiveAvgPool1d(1)` to collapse the temporal dimension into a vector of size 128. This is fed into a final `Linear` layer configured to output `num_classes` independent logits.

5.3 Training Loop (train.py)

The execution script establishes the complete PyTorch training ecosystem:

- **DataLoaders:** Converts the numpy arrays into `TensorDataset` and wraps them in `DataLoader` instances for efficient batching.
- **Imbalance Handling:** Uses `np.bincount` to dynamically calculate the distribution of the multiple classes. It passes these inverse frequencies as `weight` to the `CrossEntropyLoss` criterion, forcing the optimizer to treat rare shapes with the same importance as the overwhelming normal background noise.
- **Optimization & Validation:** Employs the Adam optimizer. A custom early stopping mechanism monitors the validation loss, saving the `best_1d_cnn_pytorch.pth` weights only when the validation loss improves.
- **Evaluation:** Evaluates the model using `torch.argmax` to determine the highest confidence class, achieving over 96% test accuracy on the multi-class simulated data.

5.4 Inference Visualization (visualize_inference.py)

To interpret the model's behavior on streaming data, a visualization script runs the trained model over a continuous segment of the dataset. It applies a rolling average to the raw probabilities to smooth out high-frequency fluctuations caused by the sliding window stride. The output is an interactive dual-axis Plotly graph rendering distinct, colored probability curves for each specific anomaly class.




Figure 4: Multi-Class 1D-CNN Confidence Probability mapped against Simulated Gamma Noise